



VASANTDADA PATIL PRATISHTHAN'S COLLEGE OF ENGINEERING AND VISUAL ARTS

ISO 9001:2015 Certified Institute

Department of Information Technology

NBA Accredited Course (Dated 01/07/2024 to 30/06/2027)

EXPERIMENT - 4

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

Theory:

After successfully creating the Kubernetes cluster, you can now deploy your application to it using the kubectl.

You need to create manifest's for your project.

- Deployment.yaml : You actual app deploying
- service.yaml : To access your app outside the network in public.

Before moving forward you can install kubectl in your system using package manager:

□ ~ □ sudo pacman -S kubectl

Some common terms in manifest:

In Kubernetes manifests, you'll encounter several common terms and concepts. Here's a brief overview of some key ones:

- kind: Should be Deployment.
- apiVersion: Specifies the API version, e.g., apps/v1.
- metadata: Contains the name, namespace, labels, and annotations of the deployment.
- spec: The specification for the deployment.
- replicas: Number of pod replicas to run.
- selector: Defines how to select pods managed by the deployment.
 - matchLabels: Key-value pairs for selecting pods.
- template: The pod template that describes the pods.
 - metadata: Labels for the pods.
 - spec: Defines the containers and their configurations.
 - containers: List of container specs.
 - name: Name of the container.

- image: Container image to use.
- ports: List of ports to expose.
- env: Environment variables for the container.
- volumeMounts: Specifies where to mount volumes in the container.
-

Service manifest:

- apiVersion: Typically v1.
- kind: Should be Service.
- metadata: Contains the name, namespace, labels, and annotations of the service.
- spec: The specification for the service.
-
- selector: Defines which pods the service will target based on labels.
- ports: List of ports exposed by the service.
 - port: Port that the service will expose.
 - targetPort: Port on the container to forward traffic to.
- type: Defines the service type (e.g., ClusterIP, NodePort, LoadBalancer).

Here's mine:

deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: node-deployment
  namespace: thrifty
  labels:
    app: node-api
spec:
  replicas: 2
  selector:
    matchLabels:
      app: node-api
  template:
    metadata:
      labels:
        app: node-api
    spec:
      containers:
        - name: node-api
          image: darkkernel/node-api
          ports:
            - containerPort: 8080
```

Service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: external-svc
  namespace: thrifty
  labels:
    app: external-svc
spec:
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 8080
      protocol: TCP
  selector:
    app: node-api
```

1. Now let's start with deployment, check if the cluster is ready.

❑ ~ ❑ kubectl get nodes

NAME

NAME	STATUS	ROLES	AGE	VERSION	
ip-192-168-26-47.ec2.internal				v1.30.2-eks-	
ip-192-168-33-0.ec2.internal	Ready	<none>	5m32s	1552ad0	v1.30.2-eks-
	Ready	<none>	5m43s	1552ad0	

2. Get your project in your system.

> git clone https://github.com/Dark-Kernel/node-api.git

Cloning into 'node-api'//.

remote: Enumerating objects: 1513, done.

remote: Counting objects: 100% (337/337), done.

remote: Compressing objects: 100% (194/194), done.

remote: Total 1513 (delta 126), reused 337 (delta 126), pack-reused 1176 (from 1)

Receiving objects: 100% (1513/1513), 2.16 MiB | 5.71 MiB/s, done.

Resolving deltas: 100% (357/357), done.

>

> cd node-api/

> ls Kubernetes/

❑ deployment.yaml ❑ service.yaml

>

3. Create Namespace if required

> kubectl create namespace thrifty

namespace/thrifty created

4. Now, create the deployment using the kubectl command.

```
> kubectl apply -f Kubernetes/deployment.yaml
deployment.apps/node-deployment created
```

5. You can check if it is applied.

```
> kubectl get deployments -n thrifty
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
node-deployment	2/2	2	2	27s

6. Then deploy the service


```
> kubectl apply -f Kubernetes/external-svc.yaml
service/external-svc created
```

7. Now, you can check your service if it is done.

```
> kubectl get svc -n thrifty
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP
external-svc	LoadBalancer	10.100.232.57	ab920d95efa3948538adb0e2601628be-396837232.us-east-1.elb.amazonaws.com

Now using the external IP, which is a subdomain actually you can access your application. And our application deployment is successful



Conclusion: Thus, we have successfully installed Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deployed our First Kubernetes Application.